

CSE202

1 - Divide and Conquer

TA: Joseph L.

I. General Concepts

- **Divide-and-conquer:** A class of algorithms in which one recursively splits sub-problems into those of smaller input size, and then later recombines the sub-problems' solutions into a solution for the original problem.
 - Splitting should **end on a base case** – a fixed-sized input for which the solution can be easily computed.
 - A complete, high-level description of the algorithm should describe **the base case, means of splitting, means of recombination, and time complexity.**
- **Master Theorem:** A theorem that quickly derives the complexity for a recurrence of the form $T(n) \leq aT\left(\frac{n}{b}\right) + f(n)$, where $f(n) = O(n^c)$ and $T(1) = 1$. This form is common among divide-and-conquer algorithms.

↳ polynomial




Figure 1. A general recurrence tree.

- Key log property: $a^{\log_b n} = n^{\log_b a}$
- The term $f(n)$ can be thought of as the complexity required to *combine subproblems*.
- This theorem involves **three** cases, depending on the relative values of $\log_b a$ and c .
 - Case 1: $c < \log_b a$
 - **The work done is bottom-heavy.** The lower layers are of higher polynomial degree than the higher ones, and thus dominate the asymptotic runtime.
 - E.g., the last layer has $O(n^{\log_b a})$, and thus has a higher degree than the top layer with $O(n^c)$ since $\log_b a > c$.
 - Therefore, **$T(n) = O(n^{\log_b a})$.**
 - Case 2: $c = \log_b a$
 - **The work done is distributed evenly across layers.** Each layer has complexity $O(n^c)$ since a given layer j has a^j subproblems, each of complexity $O\left(\left(\frac{n}{b^j}\right)^c\right)$. To elaborate:
 - $c = \log_b a$ can be rewritten as $b^c = a$.
 - Substitute: $a^j \left(\frac{n}{b^j}\right)^c = (b^c)^j \left(\frac{1}{(b^c)^j}\right) n^c = n^c$.

- There are $O(\log_b n)$ such layers.
- Therefore, $T(n) = n^c \log_b n$.
- Case 3: $c > \log_b a$
 - The work done is *top-heavy*. The higher layers are of higher polynomial degree than the lower ones, and thus dominate the asymptotic runtime.
 - E.g., the first layer has $O(n^c)$, which has higher polynomial degree than the last layer with $O(\log_b a)$, since $c > \log_b a$.
 - Therefore, $T(n) = O(n^c)$.
- Takeaways:
 - The Master Theorem is useful for quickly figuring out the complexity for a recurrence of the above form.
 - You may use the Master Theorem, but be sure to understand how each case is derived.
- Proofs of Correctness
 - The piecewise nature of divide-and-conquer lends itself naturally to **proof by induction**. Recall that a mathematical induction consists of the following parts:
 - **Base Case:** Prove that your algorithm works correctly for some basic case m_0 . Inputs falling under the base case should have sizes at or below a *small constant* – e.g., leaf nodes in the above recurrence tree.
 - **Induction Hypothesis:** Assume your algorithm works correctly for *all* inputs up to an arbitrary step \hat{m} – e.g., nodes of the \hat{m} th layer of the above recurrence tree, counting from the bottom.
 - **Induction Step:** Prove that, assuming your induction hypothesis holds, your algorithm correctly performs the next step $\hat{m} + 1$ – e.g., the nodes of layer $\hat{m} + 1$ in the above recurrence tree.
 - **Conclusion:** Formally state that **because your base case, induction hypothesis, and induction step hold, your algorithm is wholly correct for some defined set of inputs.** In our ongoing example, our base case was m_0 and we worked our way up, so **our algorithm works for every $m \geq m_0$ – i.e., every layer of our recurrence tree.**
 - Other methods of proof may be useful for proving portions of your induction (namely, your induction step).

II. Mergesort

- *Description:* We would like to efficiently sort a list of size n .
- *Splitting:* For a given list k , split into two lists i, j of equal size.
- *Base Case:* A list of size $n \leq 2$ can be sorted through simple comparison, in $O(1)$ time.
- *Combining:* Simply merge the two sorted sub-lists, in $O(n)$ time, using the merging algorithm detailed in section 2 of K&T.
- *Complexity:* $T(n) \leq 2T\left(\frac{n}{2}\right) + O(n)$
 - Master Theorem, Case 2. Complexity is **$O(n \log n)$** .
- *Correctness, by Induction*
 - *Base Case:* Sub-problems at leaf nodes have lists of size 1 or 2, and are thus trivially, correctly sorted.
 - *Induction Hypothesis:* Suppose that all sublists, **up to an arbitrary size m , are correctly sorted.** Let i, j be two such sublists.
 - *Induction Step:* We would like to prove that, if the hypothesis holds, **the combined list k (of size approximately $2m$) is correctly sorted for i, j .**
 - The merging algorithm works correctly for sorted lists, as proven in the book.
 - By the induction hypothesis, i, j are indeed sorted, thus satisfying the above precondition. Thus, k is sorted.

- *Conclusion*: Our base case, hypothesis, and step hold. Therefore, by induction, our algorithm works for lists of all sizes $n \geq 0$.
- Notes
 - $O(n \log n)$ is the best possible runtime for comparison-based sorting. This is important to remember for this class, as many problems will involve sorting as a subroutine or precomputation.
 - In the induction step, we cited without proof an algorithm that was proven in the book. In this class, you are allowed to re-use a previously discussed algorithm without proof if your solution uses it directly. If you instead use a variant of a discussed algorithm, then you must prove it from scratch.
- Related Problems: Counting inversions

III. Finding the closest pair of points on a 2D plane

- *Description*: Given n points on a plane, we would like to find the pair with the smallest distance between the two points.
- Assume, without loss of generality, that no two points share the same x or y coordinate. Whenever this is not satisfied, the coordinate plane can be rotated until it is satisfied.
- *Notation*:
 - Set of points $P = \{p_1, \dots, p_n\}$
 - $d(p_i, p_j)$ is the Euclidean distance between $p_i, p_j \in P$.
- *Preprocessing*:
 - Let P_x be P , sorted by x -coordinates.
 - Let P_y be P , sorted by y -coordinates.
 - This preprocessing can be done in $O(n \log n)$ using mergesort.
- *Splitting*: For a given set of points P , split into two sets l, r of equal size using a vertical line.




Figure 2. Points are split by a vertical line.

- *Base case*: For a set of ≤ 3 points, the closest pair can be found in constant time through pairwise comparison.
- *Combining*: For a set $S = l \cup r$, the closest pair is one of the following:
 - The closest pair in l . Let us refer to this as (q_0^*, q_1^*) . It goes without saying that since S contains all points from l , this pair is a valid candidate for the closest pair in S .
 - Analogously, the closest pair in r . Let us refer to this as (r_0^*, r_1^*) .
 - A pair containing one point from each set. We would like to determine this efficiently.
 - Let $\delta = \min(d(q_0^*, q_1^*), d(r_0^*, r_1^*))$. We are interested in whether there exist pairs closer than δ , with one point from each set.
 - Let us consider a 2δ -by- 2δ region containing at least one set of candidate points, as in Figure 4.




Figure 3. A 2δ -by- 2δ region, split into boxes of dimensions $\frac{\delta}{2}$ -by- $\frac{\delta}{2}$. The left half contains points in l , and the right half contains points in r .

- We first argue that **each of the above boxes can contain at most one point.**
 - Suppose there exists a box containing two points. Then, the distance between this pair of points is at most $\frac{\delta}{\sqrt{2}}$ (by trigonometry, if the points are on opposite corners of the box), which is smaller than δ . This contradicts our definition of δ as the *smallest distance between points on the same side*. Thus, by proof by contradiction, there cannot exist a box with more than one point.
- Let S_y be the points in S , sorted by y -coordinates. This can be derived from the pre-computed P_y in $O(n)$ time. We argue that any pair of points (p_l, p_r) s.t. $p_l \in l$ and $p_r \in r$ and $d(p_l, p_r) < \delta$, is **at most 15 positions apart in S_y .**
 - Suppose there exists a pair that is more than 15 positions apart, with $d(p_l, p_r) < \delta$. Because we established that each box has at most one point, this requires p_l, p_r to be more than 3 rows of boxes apart. This in turn means $d(p_l, p_r) > 3\left(\frac{\delta}{2}\right)$, which is larger than δ , thus violating our condition that $d(p_l, p_r) < \delta$. Thus, by proof by contradiction, p_l, p_r are at most 15 positions apart in S_y .
- Thus, each point in S has **≤ 15 candidates with which it may form the closest pair – the number is bounded by a constant**, allowing us to check all candidates in **$O(1)$ time**. Because our algorithm is exhaustive, it is guaranteed to find the closest pair, if it exists.
- We check the ≤ 15 pair candidates of up to $O(n)$ points. Each set of checks is $O(1)$. Thus, our complexity for combining is $O(n)$.
- **Complexity:** $T(n) \leq 2T\left(\frac{n}{2}\right) + O(n)$
 - Master Theorem, Case 2. The complexity is **$O(n \log n)$** .
 - The preprocessing is also $O(n \log n)$, so our total complexity is the sum: **$O(n \log n)$** .
- **Correctness by induction**
 - As is often the case with divide-and-conquer, we've already proven, in our description of the algorithm, most pieces of our proof by induction. Our proof will thus largely restate this prior work.
 - **Base Case:** For a set of ≤ 3 points, we use pairwise comparison. As this exhaustively tries all possible pairs, it is clearly correct.
 - **Induction Hypothesis:** Suppose that for all partitions smaller than or as large as some partition l, r , we have successfully computed the closest pairs.

- *Induction Step*: We would like to prove that, if our hypothesis holds, our algorithm correctly computes the closest pair for $l \cup r$. We have already proven the correctness of our “combining” step above, so we simply restate that fact.
- *Conclusion*: Our base case, hypothesis, and step hold. Therefore, our algorithm is correct for all sets of points of size $n \geq 2$ (the problem is about finding a pair, and is thus ill-posed for $n < 2$).
- Notes
 - When finding the closest pair spanning l, r , we can save some effort by eliminating all points that fall a distance $\geq \delta$ from the splitting line. This does not change the asymptotic runtime, but does eliminate unneeded effort.
 - This problem generalizes naturally to higher dimensions. In general, higher-dimensional variants are not efficiently solvable with divide-and-conquer.

IV. Integer Multiplication

- *Description*: We would like to multiply two numbers, each represented with up to n bits, efficiently.
- Introduction
 - Let x, y be two n -bit numbers – e.g., $n \geq \log_2 x$ and $n \geq \log_2 y$.
 - The elementary school algorithm for multiplying x, y is quadratic w.r.t. the number of digits.
 - Recall that you do one single-digit multiplication per pair of digits.
 - The same algorithm can be applied in binary, where it is quadratic with respect to the number of bits – i.e., it is $O(n^2)$.
 - As it turns out, there exists a divide-and-conquer algorithm of lower polynomial runtime.
- *Splitting*: We can rewrite x and y in terms of their higher and lower bits.
 - Let x_1 be the $\frac{n}{2}$ higher bits and x_0 be the $\frac{n}{2}$ lower bits. We can now rewrite x as follows:

$$x = x_1 * 2^{\frac{n}{2}} + x_0$$
 - Rewrite y analogously.
 - Now, we can rewrite the product xy as follows:

$$xy = \left(x_1 * 2^{\frac{n}{2}} + x_0\right) \left(y_1 * 2^{\frac{n}{2}} + y_0\right) = x_1 y_1 * 2^n + (x_1 y_0 + x_0 y_1) * 2^{\frac{n}{2}} + x_0 y_0$$
 - The rightmost expression seemingly has four products, of $\frac{n}{2}$ bits each, that must be computed:
 - $x_1 y_1$
 - $x_1 y_0$
 - $x_0 y_1$
 - $x_0 y_0$

$T(n) = 4T(\frac{n}{2}) + O(n)$
 - Note, however, that we do not directly need the values of $x_1 y_0$ and $x_0 y_1$ – we only need $x_1 y_0 + x_0 y_1$. We will derive this value based on $x_1 y_1$ and $x_0 y_0$ (which we do need), as well as one other multiplication, for a total of three.
 - Write: $(x_1 + x_0)(y_1 + y_0) = x_1 y_1 + x_1 y_0 + x_0 y_1 + x_0 y_0$.
 - Rewrite $x_1 y_0 + x_0 y_1$ as follows:

$$x_1 y_0 + x_0 y_1 = (x_1 + x_0)(y_1 + y_0) - (x_1 y_1 + x_0 y_0)$$
 - Now, we need only 3 products and thus 3 subproblems:
 - $(x_1 + x_0)(y_1 + y_0)$ addition is cheaper than multiplication
 - $x_1 y_1$
 - $x_0 y_0$
- *Base Case*: For $n = 1$, the numbers can simply be multiplied together.
- *Combining*: Given the products computed in our subproblems, we simply need to apply 2 bit-shifts in $O(n)$ time and add 3 numbers together, also in $O(n)$ time.

- **Complexity:** $T(n) \leq 3T\left(\frac{n}{2}\right) + O(n)$ $\log_b a = \log_2 3 > 1$
 - Master Theorem, case 1. The complexity is $T(n) = O(n^{\log_2 3}) \approx O(n^{1.59})$.
- **Correctness by induction**
 - Again, we've already proven most of what's needed here.
 - **Base Case:** Multiplication of numbers with a bounded number of bits (e.g., 1) can be done correctly and in constant time.
 - **Induction Hypothesis:** Suppose all products are correctly computed up to an arbitrary bit size m (where m is a power of 2, due to how our splitting/recombining works).
 - **Induction Step:** We would like to prove that, if our hypothesis holds, then our algorithm correctly combines the solutions of size m bits into solutions of size $2m$ bits. As the mathematics in our "splitting" and "combining" sections are algebraically sound, we simply restate that fact.
 - **Conclusion:** Our base case, hypothesis, and step hold. Therefore, our algorithm works correctly for all numbers sized $n \geq 1$ bits.

V. Convolution with Fast Fourier Transforms

- **Description:** We would like to efficiently convolve two vectors of length n .
- **Background**
 - **Convolution** of vectors a, b
 - Inputs: Vectors a of length m , b of length n
 - Output: Vector c of length $m + n - 1$.
 - $c = a * b$
 - $c_k = \sum_{i,j:i+j=k} a_i b_j$, where $0 \leq i \leq m$ and $0 \leq j \leq n$.
 - An application: multiplying polynomials $A(x)$ and $B(x)$.
 - Write $A(x)$ and $B(x)$ as vectors of coefficients:
 - $A(x) = a_0 + a_1x + \dots + a_{m-1}x^{m-1}$ as $a = (a_0, a_1, \dots, a_{m-1})$
 - $B(x) = b_0 + b_1x + \dots + b_{n-1}x^{n-1}$ as $b = (b_0, b_1, \dots, b_{n-1})$
 - If $C(x) = A(x)B(x)$, then $c = a * b$.
 - For simplicity in this section, assume $m = n$. This is without loss of generality since one can always pad the smaller vector with 0's.
 - A naïve implementation of convolution achieves complexity $O(n^2)$ by multiplying all a_i, b_j pairwise. We will discuss an $O(n \log n)$ algorithm.
 - **Complex k th roots of unity**
 - Recall the unit circle on the complex plane:

$$\begin{aligned}
 C &= A(x)B(x) \\
 &= (a_0 + a_1x + \dots + a_{m-1}x^{m-1})(b_0 + b_1x + \dots + b_{m-1}x^{m-1}) \\
 &= a_0b_0 + a_1b_0x + \dots + a_{m-1}b_0x^{m-1} + a_0b_1x + a_1b_0x^2 + \dots + a_{m-1}b_1x^m \\
 &\quad + a_0b_{m-1}x^{m-1} + a_1b_{m-1}x^m + \dots + a_{m-1}b_{m-1}x^{2m-2} \\
 &= a_0b_0 + (a_1b_0 + a_0b_1)x + \dots \\
 C_i &= \sum_{0 \leq j \leq i} a_j b_{i-j}
 \end{aligned}$$




Figure 4. The unit circle on the complex plane.

- A point on this unit circle can be written in polar coordinates as $e^{i\theta}$, where $\theta \in [0, 2\pi)$.
- The complex k th roots of unity are the k equally-spaced (in terms of angle) points $e^{2\pi ji/k}$, for $j \in \{0, 1, \dots, k-1\}$, that satisfy $(e^{i\theta})^k = 1$.




Figure 5. The complex 6th roots of unity are shown in red.

- **Polynomial interpolation**
 - Recall that a degree- d polynomial can be recovered from interpolating $d + 1$ points that lie on it.
 - This procedure takes $O(d)$ time.
- Finally, onto the algorithm...
- Key ideas
 - $C(x) = A(x)B(x)$ can be recovered from interpolating $\sim 2n$ points that lie on $A(x)B(x)$.
 - $C(x_j) = A(x_j)B(x_j)$ for any x_j . Thus, we would like to evaluate both A and B on $\sim 2n$ points, and then multiply their values together.
 - Evaluating $A(x)$ on one point takes $O(n)$ time (it's essentially a vector dot product). Naively evaluating on $2n$ points would therefore be $O(n^2)$, and thus yield no improvement.
 - **Key point:** A $\frac{k}{2}$ -th root of unity is also a k th root of unity. Thus, if we evaluate our polynomials on the k th roots of unity, we can reuse our work for the $2k$ th roots of unity. We now see the potential for a divide-and-conquer strategy.

- **Splitting**
 - The remaining sections will cover how to evaluate the polynomial A on the points; an identical procedure is also applied to B .
 - Split $A(x)$ into its odd and even coefficients:
 - $A_{\text{even}}(\hat{x}) = a_0 + a_2\hat{x} + a_4\hat{x}^2 + \dots + a_{n-2}\hat{x}^{(n-2)/2}$
 - $A_{\text{odd}}(\hat{x}) = a_1 + a_3\hat{x} + a_5\hat{x}^2 + \dots + a_{n-1}\hat{x}^{(n-2)/2}$
 - $A(x) = A_{\text{even}}(x^2) + xA_{\text{odd}}(x^2)$
 - We now have $A_{\text{even}}(\hat{x})$ and $A_{\text{odd}}(\hat{x})$ - two polynomials of halved degree, and thus two subproblems.
 - For a given subproblem, we will evaluate A_{even} and A_{odd} on the n th roots of unity.
- **Base Case:** We can compute $A(x)$ directly on the 2nd roots of unity.
- **Combining:** Given the n th roots of unity, we need to evaluate $A(x)$ on the $(2n)$ th roots of unity. Let $\omega_{j,2n} = e^{2\pi ji/2n}$ be one such root.
 - $\left(e^{\frac{2\pi ji}{2n}}\right)^2 = e^{2\pi ji/n}$, with is an n th root of unity.
 - $A(\omega_{j,2n}) = A_{\text{even}}(\omega_{j,2n}^2) + \omega_{j,2n}A_{\text{odd}}(\omega_{j,2n}^2) = A_{\text{even}}(\omega_{j,n}) + \omega_{j,n}A_{\text{odd}}(\omega_{j,n})$
 - Both terms in the rightmost expression were computed in the subproblems. Therefore, it is $O(1)$ to evaluate $A(\omega_{j,2n})$ for one value of j , and $O(n)$ for all. Combining is $O(n)$.
- **Complexity:** $T(n) \leq 2T\left(\frac{n}{2}\right) + O(n)$ a=2, b=2, C=1
 - Master Theorem, Case 2: complexity is $O(n \log n)$.
 - We perform the exact same steps for B : we spend another additive $O(n \log n)$.
 - Computing C at the end is an additive $O(n)$.
 - Total: $O(n \log n + n \log n + n) = O(n \log n)$
- **Correctness (of the evaluations on A, B) by induction**
 - **Base Case:** Our base case is a direct polynomial evaluation, and is thus correct.
 - **Induction Hypothesis:** Suppose $A_{\text{even}}, A_{\text{odd}}$ are correctly computed up to some degree n .
 - **Induction Step:** Prove that $A(x)$ is correctly computed, given the hypothesis, for the $(2n)$ th roots of unity.
 - $A(x) = A_{\text{even}}(x^2) + xA_{\text{odd}}(x^2)$ is algebraically sound.
 - $(2n)$ th roots of unity, when squared, produce n th roots of unity. Therefore, our logic for reusing past computations is algebraically sound.
 - **Conclusion:** Our hypothesis and step are correct, so by induction, our algorithm is correct for all possible values of n .
- **Notes**
 - One useful skill is the ability to recognize problems that can be modeled as convolutions.

VI. Sample Problem

We will do the following problem from Chapter 5 of K&T [1], condensed and paraphrased below.

Description: There are n bank cards and, while we can't read their numbers directly, we can compare two cards in $O(1)$ time. Among these n cards, is there a set of more than $\frac{n}{2}$ of them that are all equivalent to another? Solve this problem in $O(n \log n)$ time.

Input: n bank cards

Output: A set of $> \frac{n}{2}$ equivalent bank cards (if any)

Observations: We can compare cards pair-wise, which suggests that our method of splitting should perhaps be by halves. If we can get a recurrence tree with $O(\log n)$ layers, with each one doing a total of $O(n)$ work, then we can achieve our desired runtime of $O(n \log n)$.

Splitting: Recursively split our list of cards into halves. Each subproblem will consist of the following information:

- Which card, if any, makes up more than half of its list? This is the “solution” of our subproblem.
- Exactly how many copies of that card exist in the list?

Base Case: For $n \leq 3$ cards, we can do a fixed number of comparisons to answer both of the above questions.

Combining: Given two halves l and r and their respective solutions s_l and s_r (which can both be null), what cards are candidates for the solution s_p of the combined list p ? We argue that the only options are s_l and s_r .

- **Claim:** The solution of a combined problem p (of size n), if it exists, comprises more than half of at least one of l or r - that is, there are more than $\frac{n}{4}$ instances of it in either list.
- *Proof:* Suppose neither list has more than $\frac{n}{4}$ instances of s_p , the solution of p . Then, combined, p can have at most $\frac{n}{2}$ instances of s_p . However, we only consider a card to be a solution if it comprises more than half of the list; thus, s_p wouldn't be a solution in this case. This contradicts our knowledge that s_p is a solution and thus, by proof by contradiction, we now know that at least one of l, r must contain more than $\frac{n}{4}$ instances of s_p .

Thus, because s_l and s_r are the only cards that comprise more than half of their respective lists, they are indeed the only candidates for s_p . With this knowledge, we simply figure out, in $O(n)$ time, which (if either) of them comprises more than half of p . Note that if $s_l = s_r$, then $s_l = s_r = s_p$ without any further inspection.

Complexity: $T(n) \leq 2T\left(\frac{n}{2}\right) + O(n)$.

- *Master Theorem*, Case 2. The complexity is the desired $O(n \log n)$.

Correctness: Once again, with induction

Base Case: For $n \leq 3$ cards, we compare pair-wise, so that is trivially correct.

Induction Hypothesis: Suppose we have correctly computed solutions for subproblems of size up to n .

Induction Step: We would like to prove that if the hypothesis is true, we can correctly solve subproblems of size up to $2n$. We did this in the “Combining” section, so we refer to that.

Conclusion: Our hypothesis and step are correct so, by induction, our algorithm is correct for all possible values of n .

References

1. Kleinberg, Jon, and Éva Tardos. *Algorithm Design*. Boston: Pearson/Addison-Wesley, 2006. Print.